

MindOps 1.0

Mindwalker's AIOps Intelligence Platform for APM Managed

Enterprise-ready anomaly detection, causal RCA, and unified problem correlation with predictable performance

MindOps is an AIOps intelligence platform that extends APM Managed with advanced AI capabilities, deployed entirely inside your private network. It combines adaptive baselining, causal root cause analysis, and unified problem correlation within a single, optimized pipeline. By orchestrating anomaly detection, topology-aware fault-tree traversal, and private LLM inference, MindOps delivers precise RCA findings and enriched event data directly back into APM Managed. Designed for enterprise workloads, it guarantees sub-minute alerting, high-availability operations, and strict data sovereignty. Its architecture leverages stream processing, graph analytics, and secure local inference, ensuring consistent, production-ready insights across diverse IT environments.



Adaptive Baselining

Learns normal behavior for every metric automatically, including daily and weekly patterns, so you don't need manual thresholds.



Causal Root Cause Analysis

Uses your system's topology to find the exact root cause of a problem instead of just listing symptoms.



Problem Correlation

Groups related events into one record, preventing alert floods during cascading failures.



Private Chatbot (RAG)

A natural language assistant powered by a local AI model, with no external API calls.



APM Write-Back


Sends enriched root cause results directly back into APM Managed, so operators see more detailed insights in familiar tools.



Ops Dashboard

Sends enriched root cause results directly back into APM Managed, so operators see more detailed insights in familiar tools.

How It Works



Data Sovereignty Guarantee: All processing, storage, and AI inference occur exclusively within your network. No telemetry, metrics, logs, or LLM queries leave your environment.

Performance Highlight



Operational Efficiency

- ✓ 50% MTTR reduction
- ✓ 30–40% cost reduction
- ✓ 60% senior engineer time savings



Improved Reliability

- ✓ < 2 min alert time
- ✓ < 30 sec RCA time
- ✓ 99.99% platform availability
- ✓ 85% RCA accuracy



Business Impact

- ✓ 40% reduction in customer impact
- ✓ 100% audit compliance
- ✓ 90% noise reduction



Data Sovereignty

- ✓ 100% on-premise processing
- ✓ Zero external data transfer
- ✓ Air-gapped deployment support

Feature List

Feature	Description
CORE AI CAPABILITIES	
Adaptive Baselining	Automatically learns normal behavior for every metric with seasonal patterns (hour-of-day, day-of-week) - no manual thresholds needed
Causal Root Cause Analysis	Traverses topology using fault-tree analysis to pinpoint single root-cause entity
Problem Correlation	Groups all events sharing the same root cause into unified Problem records
Private Chatbot (RAG)	Natural language interface powered by local LLM (Ollama + Mistral 7B/Qwen2- 7B)

Feature	Description
Multimodal Analysis	Integrates logs, traces, and metrics for comprehensive analysis
DATA COLLECTION & INGESTION	
Kubernetes Integration	Pod logs, events, and audit logs via Kubernetes API and kubelet
VM Log Collection	Syslog, journald, and application logs via Fluentd/Vector agents
Database Monitoring	Query logs, slow query logs, replication status (TiDB, MySQL, PostgreSQL)
APM Integration	Distributed traces from Jaeger, Zipkin, OpenTelemetry
Metrics Collection	Import from Prometheus with AlertManager correlation
ANOMALY DETECTION	
Volume Anomaly Detection	Z-Score with sliding window for sudden log spikes/drops
Pattern Anomaly Detection	Isolation Forest algorithm for unusual log patterns
Temporal Anomaly Detection	ARIMA/Prophet models for time-series deviations
Semantic Anomaly Detection	Embedding distance calculations for novel error messages
Multi-Stage Filtering	90-95% noise reduction through intelligent filtering
VISUALIZATION & REPORTING	
Ops Dashboard	Web-based interface with live problem feed and interactive topology map
Incident timeline	Visual representation of incident progression with key events
Causal Graph Viewer	Interactive visualization of failure propagation paths
Log Explorer	Contextual log browsing with anomaly highlighting and full-text search
Service Map	Real-time topology view with health indicators
Analytics Dashboard	MTTR trends, incident patterns, SLO tracking
INTEGRATION & AUTOMATION	
Ticketing Integration	Bidirectional with ServiceNow, PagerDuty, Jira

Feature	Description
Notification Channels	Slack, Microsoft Teams, email, webhooks
REST API	Complete API coverage for incidents, analysis, topology, search
Automated Remediation	Trigger runbooks based on identified root causes
APM Write-Back	Enriched RCA results written back to APM Managed via Events API

Technical Specifications

Specification	Essential	Premium	Enterprise
FEATURE COMPARISON			
Monitored Host	Up to 100	Up to 500	Unlimited
Data Retention	30 days	1 year	5 years
Concurrent Users	10	50	Unlimited
Support SLA	24hr response	4hr response	1hr response + Dedicated TAM
CORE AI CAPABILITIES			
Adaptive Baselining	✓	✓	✓
Causal Root Cause Analysis	✓	✓	✓
Problem Correlation	✓	✓	✓
Multi-Stage Noise Filtering	✓	✓	✓
RCA Accuracy SLA	80%	85%	90% +
DATA COLLECTION & INTEGRATION			
Kubernetes Log Collection	✓	✓	✓
VM Log Collection	✓	✓	✓

Specification	Essential	Premium	Enterprise
Database Monitoring	MySQL only	MySQL, PostgreSQL	MySQL, PostgreSQL, TiDB, Oracle
APM Integration	OpenTelemetry	OpenTelemetry, Jaeger	All (OpenTelemetry, Jaeger, Zipkin, custom)
Prometheus Metrics	✓	✓	✓
Custom Log Sources	✗	5 sources	Unlimited
AI & LLM FEATURES			
Private LLM (Ollama)	Mistral 7B	Mistral 7B / Qwen2-7B	Any 7B-70B model
CoPilot Chatbot	✗	✓	✓
RAG Knowledge Base	✗	✓ (10 GB)	✓ (Unlimited)
Natural Language Insights	Basic	Advanced	Advanced + Custom Prompts
LLM Response Time (GPU)	N/A	< 20 sec	< 10 sec
VISUALIZATION & DASHBOARD			
Ops Dashboard	✓	✓	✓
Incident Timeline	✓	✓	✓
Causal Graph Viewer	✓	✓	✓
Log Explorer	Basic	Advanced + Filters	Advanced + Full-text + Regex
Service Map	Static	Interactive	Interactive + Custom Views
Analytics Dashboard	✗	✓	✓ + Custom Reports
Custom Log Sources	1 min refresh	30 sec refresh	Real-time WebSocket
INTEGRATION & AUTOMATION			
REST API Access	Read-only	Full API	Full API + Custom Endpoints
Ticketing Integration	✗	ServiceNow, Jira	ServiceNow, PagerDuty, Jira + Custom
Notification Channels	Email	Email, Slack	Email, Slack, Teams, Webhooks

Specification	Essential	Premium	Enterprise
APM Write-Back	✗	✓	✓
Automated Remediation	✗	✗	✓ (Runbook triggers)
SSO Integration	✗	SAML	SAML, OIDC, Azure AD, Okta
SECURITY & COMPLIANCE			
TLS Encryption	✓	✓	✓
Encryption at Rest	✓	✓	✓
RBAC	2 roles	4 roles	Unlimited custom roles
Audit Logging	7 days	90 days	Unlimited + SIEM export
PII Redaction	✗	✓	✓ + Custom patterns
Air-Gapped Deployment	✗	✗	✓
Compliance Reports	✗	Basic	SOC2, ISO27001, custom
DEPLOYMENT & INFRASTRUCTURE			
Deployment Model	On-Premise	On-Premise / Hybrid	On-Premise / Hybrid / Air-Gapped
High Availability	✗	✓	✓ + Multi-AZ
Disaster Recovery	Manual	Automated (RTO 1hr)	Automated (RTO 15 min)
GPU Support	✗	Optional	Required (included)
Kubernetes Support	K3s	K3s / K8s	K3s / K8s / OpenShift
SUPPORT & SERVICES			
Documentation	✓	✓	✓
Email Support	✓	✓	✓
Phone Support	✗	Business hours	24/7
Dedicated TAM	✗	✗	✓
Onboarding Training	Self-service	2 sessions	Unlimited + Custom

Specification	Essential	Premium	Enterprise
Professional Services	X	Add-on	Included (40 hrs/year)
Platform Uptime SLA	99.5%	99.9%	99.9%

System Requirement Specifications by Tiers

	Essential	Premium	Enterprise
Ideal For	✓ Small teams ✓ Pilot deployments ✓ Development / staging environments	✓ Mid-size organization ✓ Production environments ✓ Teams requiring advanced features	✓ Large organizations ✓ Regulated industries ✓ Mission-critical deployments ✓ Multi-region setups
INFRASTRUCTURE REQUIREMENTS			
AI Processing	1 node: 16 vCPU / 64 GB RAM / 250 GB SSD	2 nodes (HA): 32 vCPU / 128 GB RAM / 500 GB SSD each	4+ nodes (HA): 64 vCPU / 256 GB RAM / 1 TB SSD each
LLM Inference	-	1 node (GPU): 16 vCPU / 64 GB RAM / 200 GB NVMe / NVIDIA A10G	2+ nodes (HA): 32 vCPU / 128 GB RAM / 500 GB NVMe / NVIDIA A100 each
Data Storage	1 node: 8 vCPU / 32 GB RAM / 1 TB SSD	3 nodes: 16 vCPU / 64 GB RAM / 4 TB HDD each	6+ nodes: 32 vCPU / 128 GB RAM / 8 TB SSD each
Topology Graph	-	2 nodes (HA): 8 vCPU / 32 GB RAM / 500 GB SSD each	3+ nodes (Cluster): 16 vCPU / 64 GB RAM / 1 TB SSD each
Kafka Broker	-	3 nodes: 8 vCPU / 32 GB RAM / 2 TB HDD each	5+ nodes: 16 vCPU / 64 GB RAM / 4 TB SSD each

	Essential	Premium	Enterprise
App Layer	1 node: 4 vCPU / 8 GB RAM / 100 GB SSD	2 nodes (HA): 8 vCPU / 16 GB RAM / 200 GB SSD each	4+ nodes (HA): 16 vCPU / 32 GB RAM / 500 GB SSD each
GPU	Not required (CPU inference only)	-	-
Utility/Ops	-	1 node: 4 vCPU / 16 GB RAM / 200 GB SSD	2 nodes (HA): 8 vCPU / 32 GB RAM / 500 GB SSD each
PERFORMANCE TARGET			
Anomaly Alert	< 5 min	< 2 min	< 1 min
Root Cause Analysis	< 2 min	< 30 sec	< 15 sec
CoPilot Response	-	< 20 sec	< 10 sec
Dashboard Query	< 5 sec P99	< 3 sec P99	< 1 sec P99
Platform Uptime	-	99.9%	99.99%

Enterprise Tier Scalability

- Log Volume: Up to 10M events/sec
- Services Monitored: Up to 10,000 services
- Concurrent Users: Unlimited

Software Pre-Requisites

Component	Version	Notes
Operating System	Ubuntu 24.04 LTS	Or Compatible Linux distribution
Container Runtime	Docker 24+ / containerd	Required for K3s
Orchestration	K3s (Kubernetes)	Lightweight self-hosted K8s
Service Mesh	Istio 1.21+	mLTS, zero-trust network policy

Network Requirements

Connection	Protocol	Purpose
MindOps → APM Managed	HTTPS:8443	Or Compatible Linux distribution
Operator → Dashboard/API	HTTPS:443	Internal network access
MindOps → ITSM Systems	HTTPS/SMTP	Alert dispatch (ServiceNow, Jira, email, Slack)
Internet/External	None	No external connectivity required or permitted

Add-On Services

Add-On	Description	Available For
Additional Monitored Hosts	Expand beyond tier limits	All tiers
Extended Data Retention	Increase retention beyond tier default	Essential, Premium
GPU Infrastructure	Add GPU nodes for CoPilot	Essential , Premium
Custom Integration Development	Build custom connectors	Premium, Enterprise
On-Site Training	Customized training sessions	All tiers
24/7 Support Upgrade	Upgrade to round-the-clock support	Essential , Premium
Air-Gapped Deployment Package	Enable fully isolated deployment	Premium
Compliance Certification Support	Assistance with audit preparation	Premium, Enterprise

Deployment Models

Tiers	On-Cloud	On-Premise (Full deployment within customer data center)	Hybrid (On-premise processing with optional cloud backup)	Air-Gapped (Completely isolated with no external connectivity)
Essential	✓	✓	✗	✗
Premium	✓	✓	✓	✗
Enterprise	✓	✓	✓	✓

Contact

+62 21 2880 2137

PT. Reka Citra Solusi

Ruko Grand Centro Bintaro Blok A19

JL. Kodam Bintaro Raya, Pesanggrahan

Jakarta Selatan

Jakarta 12320

Indonesia

Website

<https://www.mindwalker.ai>

2026-WW10-V1.0-EN